


# Peer to Peer II

## Tapestry

# Overlay Routing Concerns

- **Stretch**
  - routing delay penalty (RDP)
- **Load balancing**
  - popular object located on only one node




# What to Do?

- **Have multiple copies of objects at well distributed nodes**
  - How many?
  - If you have an object at all nodes, what is the cost of read, what is the cost of insert/delete?
  - What if you have one object?
- **Take communication distance into account when setting up overlay networks**


# Tapestry

**Data Block 1**  
0x87a6df52

I want  
Block 1




0x87a6df52 locations:  
0x2a74ca56  
0xd53b7621




- Assign each block a unique m-bit ID
  - crypto hash of its contents
- Assign each computer a unique m-bit ID
- Store multiple copies of blocks each at a number of computers
- Store block addresses at computer that has closest ID
  - addresses are cached at other nodes
- Route requests for that block to that computer
  - request is redirected to nearest computer that has copy of block

# How to Route?



# Prefix Routing



# Neighbor Table for 3312

	0	1	2	3
xxxx	<b>0331</b> 128.148.128.173	<b>1332</b> 128.138.117.92	<b>2302</b> 128.118.165. 27	<b>3312</b> 128.213.97.6
3xxx	—	<b>3111</b> 128.172.53.237	—	<b>3312</b> 128.213.97.6
33xx	—	<b>3311</b> 128.12.236.81	<b>3320</b> 128.248.192.76	—
331x	—	<b>3311</b> 128.12.236.81	<b>3312</b> 128.213.97.6	—

# Routing Algorithm

**// executed at each node in route to destination**

```
NextHop(targetHash, step) {  
    nextDigit = digit(targetHash, step)  
    return(table[step, nextDigit])  
}
```

```
digit(num, pos) {  
    return ((num/based-pos)%base)  
}
```



# Surrogate Routing

- Store object's location list at unique computer whose hash is "closest" to the object's hash
    - unique computer known as the *root*
    - all routes to object's hash reach the root regardless of the starting point
    - the path to this root goes through various "surrogate" nodes
      - if there is a hole in the neighbor table corresponding to the "next digit", then choose the first non-empty entry in the row that's greater (mod base) than the one desired
      - the node routed to is the *surrogate*
-

# How?

- If no next hop exists, try the next larger digit, mod *base*
  - each neighbor-table row must have at least one entry
    - why?
  - if any two neighbor-table rows (of different nodes) share the same prefix, they must agree on which entries are null
    - why?

# Neighbor Tables for 3312 and 3320

	0	1	2	3
xxxx	0331 128.148.128.173	1332 128.138.117.92	2302 128.118.165.27	3312 128.213.97.6
3xxx	—	3111 128.172.53.237	—	3312 128.213.97.6
33xx	—	3312 128.213.97.6	3320 128.248.192.76	—
331x	—	3311 128.12.236.81	3312 128.213.97.6	—

	0	1	2	3
xxxx	0121 128.148.158.13	1001 128.18.11.192	2130 128.113.225.127	3311 128.12.236.81
3xxx	—	3120 128.162.253.247	—	3320 128.248.192.76
33xx	—	3311 128.12.236.81	3320 128.248.192.76	—
332x	3320 128.248.192.76	—	—	—

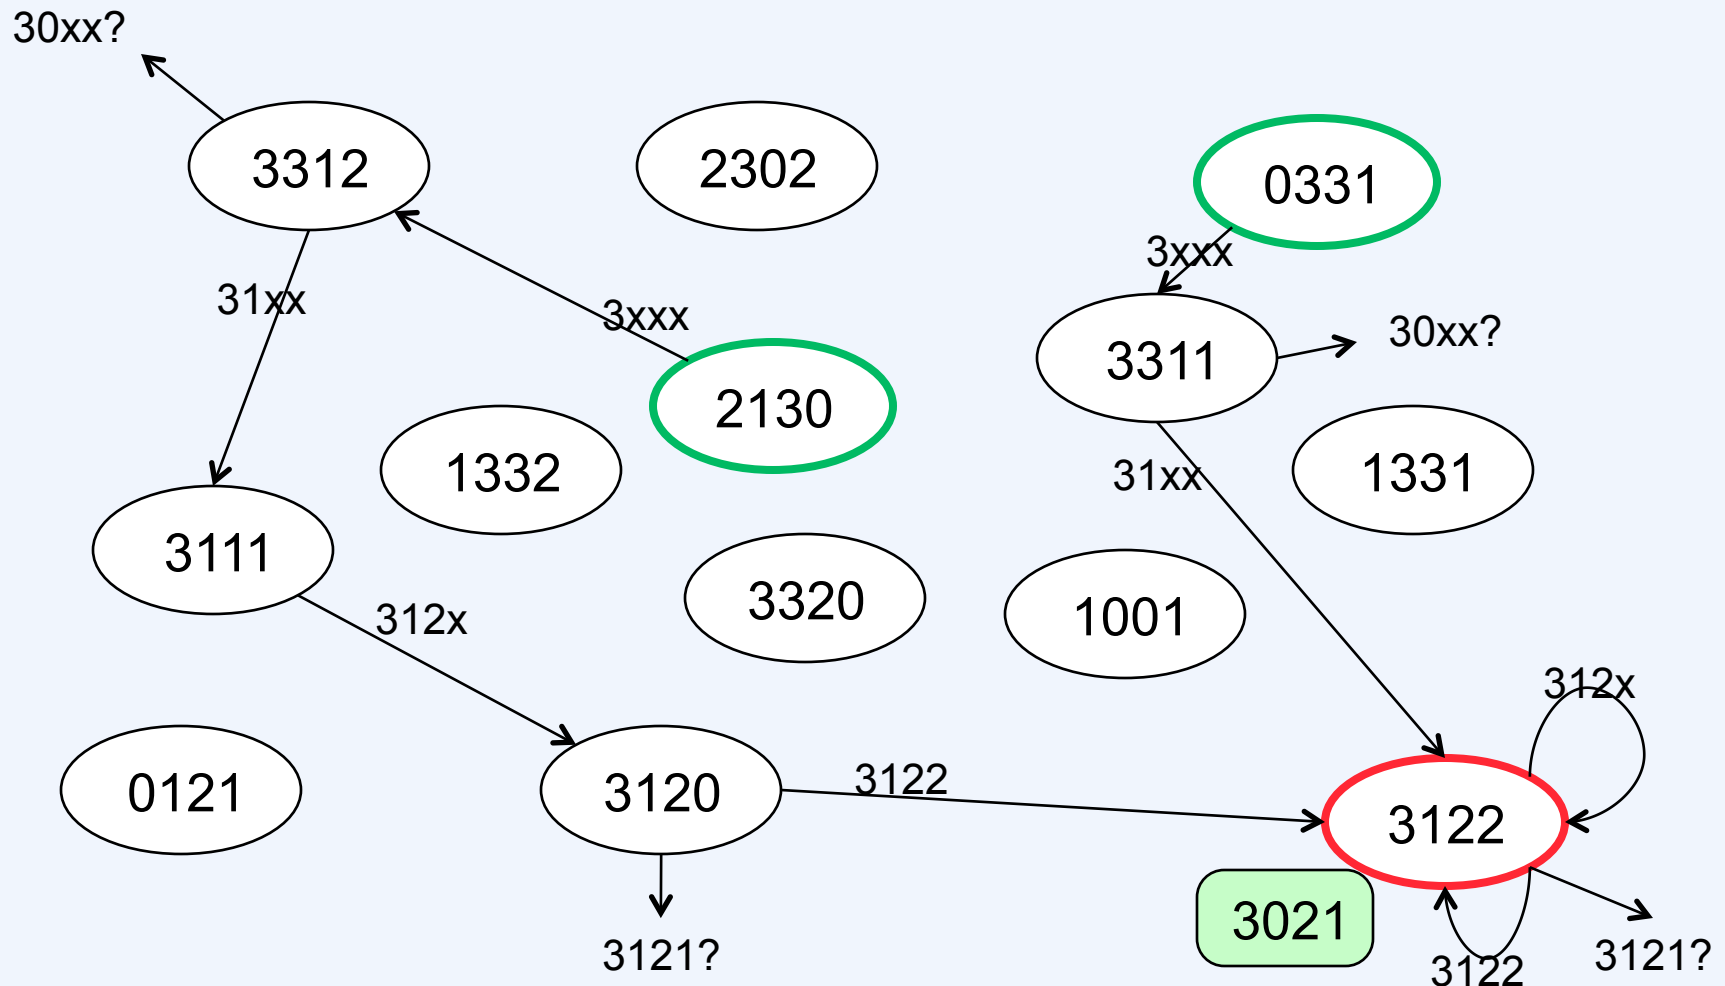
# Surrogate Routing Algorithm

// executed at each node in route to destination

```
NextHop(targetHash, step) {  
    nextDigit = digit(targetHash, step)  
    while ((next = table[step, nextDigit]) == NULL)  
        nextDigit += 1 mod base  
    return next  
}
```

```
digit(num, pos) {  
    return ((num/based-pos)%base)  
}
```


# Surrogate Routing for 3021




# Performance and Redundancy

- For any particular neighbor-table entry, there may be a number of possible valid next hops
  - all of them work
  - choose the one that's "closest"
    - communication delay makes sense for this
  - if a next hop can't be reached
    - use one of the other possible next hops
    - store some number of them in table
      - "secondary entries"

# Publishing

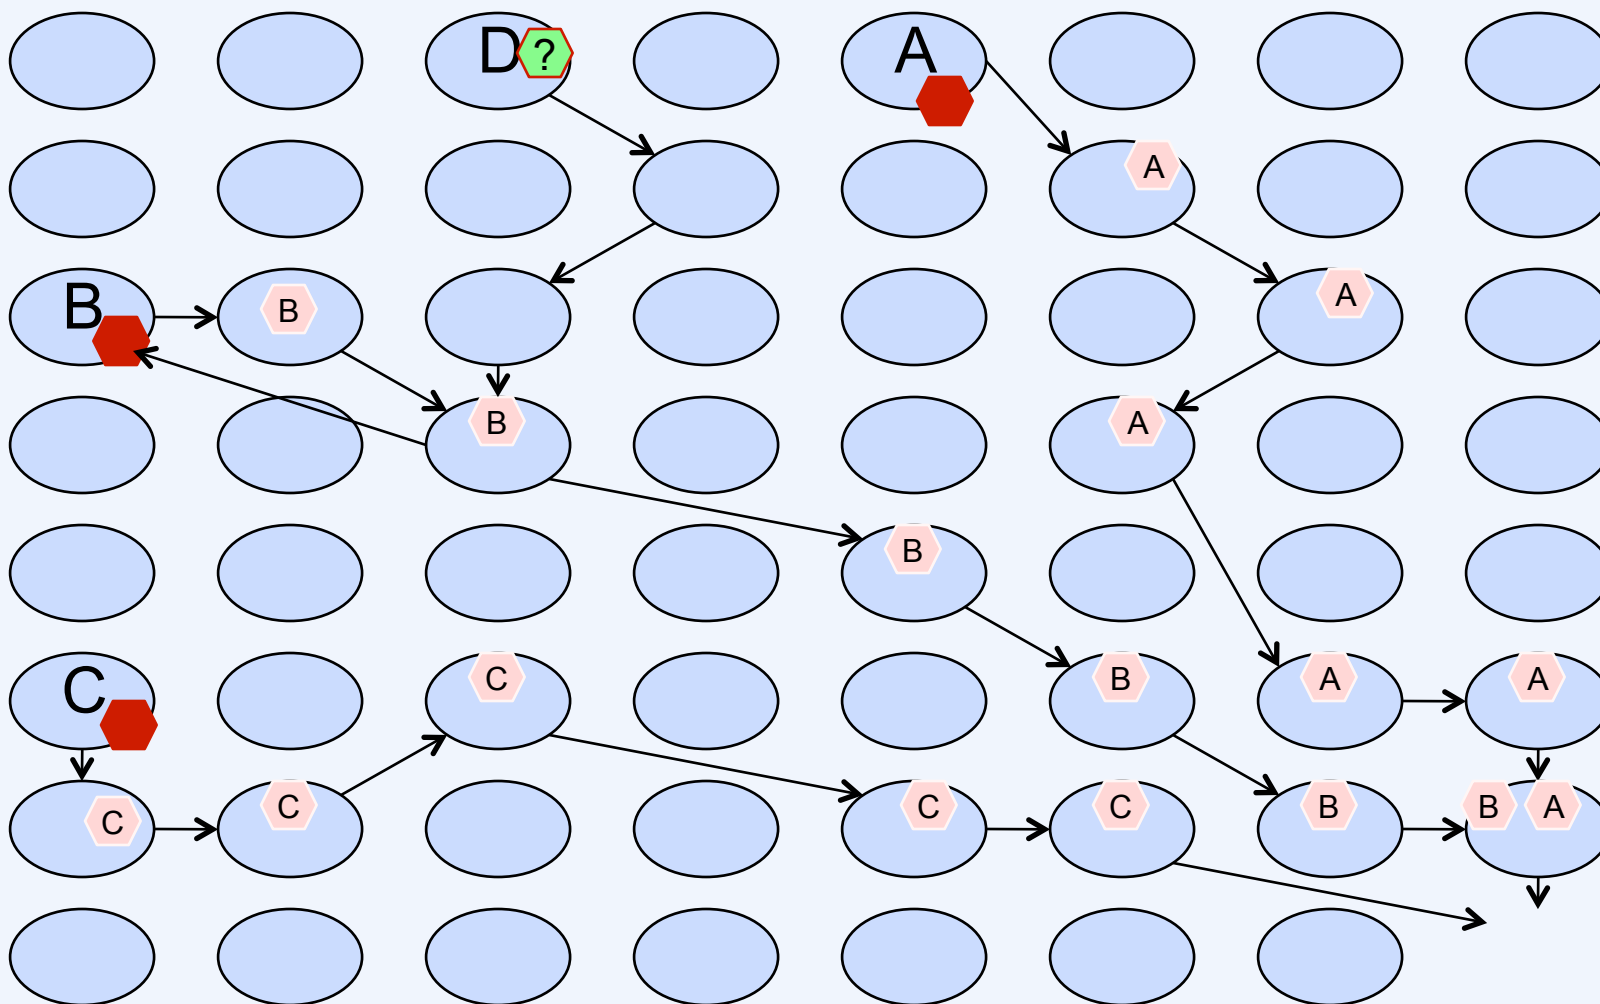


# Failure






# Failure



# Soft State


- **State information times out**
  - e.g., reference to node holding an object
- **Must be periodically reestablished**
  - nodes must periodically republish their objects

# Recovery



# Redundant Redundancy

- When “root nodes” of objects disappear, it may take some time before re-publication is effective
  - solution: extra root nodes
  - how?
    - “salt” the hashes
      - append a small integer before hashing
      - multiple hashes for one object: multiple roots for the object



# Adding a Node

- **Steps for adding node n**
  - 1) **find existing node G**
  - 2) **search for n's hash starting at G**
  - 3) **at each step i, fill in row i of n's table with row i of table of node being visited**
  - 4) **stop when empty table entry is encountered**
  - 5) **fill in remainder of table with self entries**
  - 6) **notify other nodes to update their tables**

# Initializing a Neighbor Table Row (1)

	0	1	2	3
xxxx	0331 128.148.128.173	1332 128.138.117.92	2302 128.118.165. 27	3312 128.213.97.6
3xxx	—	—	—	—
30xx	—	—	—	—
300x	—	—	—	—

n's (3001's) Table

	0	1	2	3
xxxx	0331 128.148.128.173	1332 128.138.117.92	2302 128.118.165. 27	3312 128.213.97.6
3xxx	—	3111 128.172.53.237	—	3312 128.213.97.6
33xx	—	3311 128.12.236.81	3320 128.248.192.76	—
331x	—	3311 128.12.236.81	3312 128.213.97.6	—

G's (3312's) Table

# Initializing a Neighbor Table Row (2)

	0	1	2	3
xxxx	0331 128.148.128.173	1332 128.138.117.92	2302 128.118.165.27	3312 128.213.97.6
3xxx	3001 128.250.19.172	3111 128.172.53.237	—	3312 128.213.97.6
30xx	—	—	—	—
300x	—	—	—	—

n's (3001's) Table

	0	1	2	3
xxxx	0331 128.148.128.173	1332 128.138.117.92	2302 128.118.165.27	3312 128.213.97.6
3xxx	—	3111 128.172.53.237	—	3312 128.213.97.6
33xx	—	3311 128.12.236.81	3320 128.248.192.76	—
331x	—	3311 128.12.236.81	3312 128.213.97.6	—

step 2 (3312's) Table

# Initializing a Neighbor Table Row (3)

	0	1	2	3
xxxx	0331 128.148.128.173	1332 128.138.117.92	2302 128.118.165.27	3312 128.213.97.6
3xxx	3001 128.250.19.172	3111 128.172.53.237	—	3312 128.213.97.6
30xx	3001 128.250.19.172	—	—	—
300x	—	—	—	—

n's (3001's) Table

	0	1	2	3
xxxx	0331 128.148.128.173	1332 128.138.117.92	2302 128.118.165.27	3312 128.213.97.6
3xxx	—	3111 128.172.53.237	—	3312 128.213.97.6
31xx	—	3111 128.172.53.237	3120 128.162.253.247	—
311x	—	3111 128.172.53.237	—	—

step 3 (3111's) Table



# Initializing a Neighbor Table Row (4)


	0	1	2	3
xxxx	0331 128.148.128.173	1332 128.138.117.92	2302 128.118.165.27	3312 128.213.97.6
3xxx	3001 128.250.19.172	3111 128.172.53.237	—	3312 128.213.97.6
30xx	3001 128.250.19.172	—	—	—
300x	—	3001 128.250.19.172	—	—

n's (3001's) Table

	0	1	2	3
xxxx	0331 128.148.128.173	1332 128.138.117.92	2302 128.118.165.27	3312 128.213.97.6
3xxx	—	3111 128.172.53.237	—	3312 128.213.97.6
31xx	—	3111 128.172.53.237	3120 128.162.253.247	—
311x	—	3111 128.172.53.237	—	—

step 4 (3111's) Table

# Updated View



# Notifying Others

- Need to insert  $n$  in all neighbor-table entries that are empty where  $n$  should go
- Before  $n$  was added, any search for  $n$  ended up at its root
- Proceed backwards from root
  - each neighbor table includes back pointers to all nodes that route to it
  - *flooding* procedure:
    - on receipt of notification
      - if routing table contains hole where  $n$  goes
        - insert  $n$
        - notify neighbors via back pointers

# A Problem

	0	1	2	3
xxxx	0121 128.148.158.13	1001 128.18.11.192	2130 128.113.225. 127	3311 128.12.236.81
3xxx	—	3120 128.162.253.247	—	3320 128.248.192.76
33xx	—	3311 128.12.236.81	3320 128.248.192.76	—
332x	3320 128.248.192.76	—	—	—

- Node 3322 is added
- Object 3321's surrogate was 3320
  - now it's 3322
  - what about all the location info that was stored assuming 3320?

# Doing a Better Job ...

- Find all nodes for which new node fills holes in neighbor tables
  - propagate new node on spanning tree of just the relevant nodes
- Handle “re-rooted” objects
  - efficiently ...
- Build new neighbor tables
  - optimizing for closeness

# Observation

- Let  $\alpha$  be longest common prefix of new node and its root
- All nodes whose neighbor tables contain holes to be filled by new node have this prefix

	0	1	2	3
xxxx	0331 128.148.128.173	1332 128.138.117.92	2302 128.118.165.27	3312 128.213.97.6
3xxx	3001 128.250.19.172	3111 128.172.53.237	—	3312 128.213.97.6
30xx	3001 128.250.19.172	—	—	—
300x	—	3001 128.250.19.172	—	—

n's (3001's) Table

	0	1	2	3
xxxx	0331 128.148.128.173	1332 128.138.117.92	2302 128.118.165.27	3312 128.213.97.6
3xxx	—	3111 128.172.53.237	—	3312 128.213.97.6
31xx	—	3111 128.172.53.237	3120 128.162.253.247	—
311x	—	3111 128.172.53.237	—	—

root's (3111's) Table

# Application

- **Send new node's info to all nodes with prefix  $\alpha$** 
  - hash and IP address
- **How?**
  - via spanning tree that reaches all such nodes
  - use “acknowledged multicast”

# Acknowledged Multicast

```
n.acknowledgedMulticast( $\alpha$ , function) {  
  if (notOnlyNodeWithPrefix( $\alpha$ ))  
    for i = 0 to b-1  
      neighbor = neighborWithPrefix( $\alpha \cdot i$ )  
      if neighbor exists  
        S = neighbor.acknowledgedMulticast( $\alpha \cdot i$ , function)  
    else  
      apply function  
    wait S  
    SendAcknowledgement()  
}
```

---



# Using Acknowledged Multicast

- To fill holes in neighbor tables with new node
  - supply function that does this
- To “re-root” object references (move them to the new node)
  - supply function that does this
- To get list of all nodes with given prefix
  - supply function that returns node IDs

# Races

- **Suppose a search for object  $x$  occurs while new node  $N$  is being added**
  - $N$  becomes the new root for  $x$
- **Potential problems:**
  - 1) **search arrives at  $N$  before object references are transferred to  $N$**
  - 2) **search arrives at old root after object references are transferred to  $N$**

# Problem 1 Solution

- **Mark N as “being inserted”**
  - on “object not found”
    - forward request to original root

# Problem 2 Solution

- **Include path with search request**
  - on “object not found”
    - **check neighbor table to see if path would still be taken**
      - **i.e., has hole subsequently been filled (by new root)?**
      - **if so, reroute to new root**

# Optimizing the Neighbor Table

- **Want primary node for each neighbor-table entry to be the one that's closest**
  - **secondary nodes should be the next closest**
- **How can this be constructed for a new node?**

# Sketch

- Use acknowledged multicast to find all nodes with prefix  $\alpha$  (longest prefix in common with root)
- From this set of nodes, construct table row  $|\alpha|$ 
  - determine which are closest to the new node
- From this set of nodes, find all nodes with prefix one shorter than  $\alpha$ 
  - construct table row  $|\alpha|-1$ , using closest nodes
- etc.